

Budget Tracker — Product Requirements Document

Version: 1.1

Date: May 2026

Status: Draft

Hackathon: Accelerator Build Day

Stack: Bolt (frontend) · Supabase (backend) · OpenAI API (AI layer)

1. Problem Statement

Most people have a vague, anxious relationship with their money. They know they're overspending somewhere — they just don't know where, and they don't want to sit down and figure it out. Existing tools (banking apps, spreadsheets, Monzo budgets) show raw data but offer no interpretation, no nudge, and no plain-English insight.

The gap is not the data. It's the meaning.

This app closes that gap by letting users log expenses in natural language and receive AI-generated insights that feel like a conversation with a financially savvy mentor — not a spreadsheet.

Core pain points

1. Manual categorisation of expenses is tedious and error-prone
 2. Financial dashboards show numbers but not what to do about them
 3. People don't know which of their habits are actually costing them the most
-

2. Target User

Primary persona: "Anxious-but-aware adult"

- Age 22-40, employed, living independently — globally distributed
 - Has a rough sense of their budget but no system
 - Uses their banking app reactively (checks after overspending, not before)
 - Would benefit from a weekly nudge that doesn't require effort to set up
 - Not a spreadsheet person; values simplicity and plain language
-

3. Goals & Success Metrics

Goal	Metric
Fast, frictionless expense entry	User can log an expense in under 10 seconds
Accurate AI categorisation	≥90% of entries correctly categorised without user correction
Meaningful AI insights	Dashboard surfaces at least one actionable insight per session
Global usability	Currency correctly detected or selected for all supported locales
MVP completeness	All core features working end-to-end for demo

4. Feature Set

4.1 MVP Features (hackathon scope)

F1 — Natural language expense entry

Users type expenses in plain English rather than filling in forms.

- Input: free-text field (e.g. "spent \$12 on lunch" or "¥800 on the subway")
- AI parses: amount, currency, description, date (defaults to today), category
- Currency inferred from symbol in input — falls back to user's selected currency
- User sees a confirmation card before saving
- Language: English only (v1)

F2 — AI auto-categorisation

Every expense is automatically assigned a category by the AI.

Categories: Food & Drink · Transport · Shopping · Entertainment · Health · Subscriptions · Bills · Other

- User can override the category with a single tap
- Category is stored and used in all summary views

F3 — Expense dashboard

A clean summary view showing spending at a glance.

- Total spent this week / this month (displayed in user's selected currency)
- Spending by category (bar chart or donut)
- List of recent expenses with category badge
- "Essential vs. avoidable" split (AI-generated tag per expense)

F4 — AI weekly summary ("Your Money Mentor")

A plain-English paragraph generated by AI each time the user opens the dashboard, written in the voice of an encouraging financial mentor.

Example output:

"You're making real progress this week — your total came in at \$340, and I noticed you cut back on dining out compared to last week. That matters. The one thing I want you to look at: four subscription charges totalling \$38.97. Ask yourself honestly — are all of these still serving you? Cancelling even one is a win."

- Triggered on dashboard load (cached, regenerated once per day)
- Warm, direct, empowering tone — see Brand Voice (Section 11)
- Always ends with one specific, actionable suggestion

F5 — Currency localisation

Users select their home currency on first login. All amounts display in that currency throughout the app.

- Currency selector on onboarding screen (searchable dropdown)
- Selection stored in user profile (Supabase)
- AI parsing prompt passes selected currency as context
- Amounts stored as raw numeric values; currency code stored separately
- Display formatting uses locale-appropriate separators (e.g. 1,000.00 vs 1.000,00)

Supported currencies (MVP):

Region	Currency	Symbol	Code
United States	US Dollar	\$	USD
United Kingdom	Pound Sterling	£	GBP
Eurozone	Euro	€	EUR
Japan	Japanese Yen	¥	JPY
Canada	Canadian Dollar	CA\$	CAD
Australia	Australian Dollar	A\$	AUD
India	Indian Rupee	₹	INR
Brazil	Brazilian Real	R\$	BRL
South Africa	South African Rand	R	ZAR
UAE	UAE Dirham	د.إ.	AED
Singapore	Singapore Dollar	S\$	SGD
Switzerland	Swiss Franc	CHF	CHF
Mexico	Mexican Peso	MX\$	MXN
South Korea	South Korean Won	₩	KRW
Nigeria	Nigerian Naira	₦	NGN

Currency conversion is out of scope for v1. All expenses are recorded in the user's selected currency. Multi-currency per-entry support is a v2 feature.

F6 — Saved expenses (Supabase persistence)

All expenses are saved per user session with Supabase.

- User identified by simple email login (Supabase Auth)
 - Expenses stored with: amount, currency_code, description, category, date, essential_flag, created_at
 - Data persists across sessions
-

4.2 Stretch Features (post-MVP / if time allows)

Feature	Description
Monthly budget cap	User sets a monthly limit per category; dashboard shows progress bar
Recurring expense detection	AI flags likely subscriptions after 2+ similar entries
Export to CSV	Download all expenses as a spreadsheet
Spending streak	"You've logged expenses for 5 days in a row" motivational nudge
Currency auto-detect	Infer currency from browser locale on first load

5. AI Integration Specification

5.1 Expense parsing (F1 + F2)

Trigger: User submits natural language expense entry

Model: gpt-4o-mini (fast, cheap, sufficient for parsing)

Prompt pattern:

You are a budget tracking assistant. Parse the following expense entry and return JSON only.

Entry: "{user_input}"

User's selected currency: {currency_code} ({currency_symbol})

Return:

```
{
  "amount": number (raw numeric value, no currency symbol),
  "currency_code": string (3-letter ISO code – use user's selected currency unless
a different symbol is explicit in the input),
  "description": string (clean, concise),
  "category": one of ["Food & Drink", "Transport", "Shopping", "Entertainment",
"Health", "Subscriptions", "Bills", "Other"],
  "is_essential": boolean,
  "confidence": number between 0 and 1
}
```

If amount cannot be determined, return null for amount.

Today's date is {today}.

Output handling: Parse JSON response, show confirmation card to user with formatted currency amount, allow edits before saving.

5.2 Weekly AI summary — "Your Money Mentor" (F4)

Trigger: Dashboard load (once per day, cached in Supabase)

Model: `gpt-4o-mini`

Prompt pattern:

You are an encouraging financial mentor in the style of Suze Orman – warm, direct, empowering, and honest. You believe everyone can take control of their finances with the right guidance.

Based on this user's spending data for the past 7 days, write a 3-4 sentence summary.

Spending data: {json_summary_of_expenses}

User's currency: {currency_code}

Rules:

- Open by acknowledging something positive in their spending this week
- Name one specific area of concern – be honest but never shame them
- End with one concrete, actionable suggestion framed as an empowering choice
- Tone: mentor not lecturer. Encouraging not cheerleading. Direct not cold.
- Use their currency symbol when citing amounts
- Do not use bullet points. Write in flowing sentences.

6. Data Schema (Supabase)

Table: **user_profiles** (profile extension)

Extends Supabase Auth with currency preference.

Column	Type	Notes
id	uuid	Matches auth.users id
currency_code	text	ISO 4217 code, e.g. "USD"
currency_symbol	text	Display symbol, e.g. "\$"
created_at	timestampz	Auto-set

Table: **expenses**

Column	Type	Notes
id	uuid	Primary key
user_id	uuid	Foreign key → auth.users

Column	Type	Notes
amount	numeric(14,2)	Raw numeric value
currency_code	text	ISO 4217 code at time of entry
description	text	Cleaned description
category	text	One of 8 categories
is_essential	boolean	AI-generated, user-overridable
raw_input	text	Original user input (for debugging)
ai_confidence	numeric(3,2)	0.00–1.00
expense_date	date	Date of expense (not entry)
created_at	timestamptz	Auto-set

Table: `ai_summaries`

Column	Type	Notes
id	uuid	Primary key
user_id	uuid	Foreign key → auth.users
summary_text	text	AI-generated paragraph
generated_at	timestamptz	For cache invalidation
week_start	date	Monday of the summarised week

7. User Flows

Flow 1 — Onboarding & currency selection

First visit

- Sign up with email (Supabase Auth)
- Currency selection screen (searchable list of 15 supported currencies)

- Selection saved to user_profiles
- Lands on dashboard (empty state with prompt to add first expense)

Flow 2 — Add an expense

- Open app
- Tap "Add expense"
 - Type in natural language: "grabbed a coffee, \$4.50"
 - AI parses → confirmation card shows (amount in selected currency, category, essential flag)
 - User confirms or edits
 - Saved to Supabase
 - Returns to dashboard (updated)

Flow 3 — View dashboard

- Open app
- Dashboard loads
 - Supabase fetches this week's expenses
 - AI summary generated (or served from cache)
 - User sees: total in their currency, category breakdown, recent list, AI mentor insight

Flow 4 — Override a category

- Dashboard → tap expense item
- Category picker appears
 - User selects correct category
 - Updated in Supabase
 - Dashboard recalculates

8. Out of Scope (for hackathon)

- Bank/Open Banking integration (Plaid, TrueLayer)
- Native mobile app
- Currency conversion between entries
- Multi-language UI (English only for v1)
- Shared/household budgets

- Historical data beyond 30 days
- Push notifications

9. Demo Script (for judging)

Narrative hook (30 seconds):

"Everyone knows roughly how much they earn. Almost nobody knows where it goes. Banking apps give you data — but data without interpretation is just noise. We built something different. A financial mentor in your pocket."

Live demo (90 seconds):

1. Show onboarding currency selector — pick USD, then switch to GBP to show global support
2. Type: "spent £38 on a big Sainsbury's shop" — show AI parse it correctly as Food & Drink, essential
3. Type: "Netflix £15.99" — show it flagged as Subscription, not essential
4. Type: "Uber home after the pub, £14" — show Transport, not essential
5. Open dashboard — show category breakdown and AI mentor summary
6. Read the actionable suggestion aloud — land the Suze Orman tone

Close:

"The AI does the categorisation, spots the patterns, and speaks to you like a mentor — not a bank statement. No setup. No spreadsheets. Available to anyone, anywhere in the world."

10. Build Priorities & Team Split

Area	Owner	Priority
Expense entry UI + confirmation card	Bolt / frontend	P0
Supabase schema + auth	Backend	P0
OpenAI parsing integration	AI	P0
Dashboard + category chart	Bolt / frontend	P0
Currency selector (onboarding)	Frontend + Backend	P0
AI summary generation + caching	AI	P1

Area	Owner	Priority
Essential vs. avoidable flag UI	Frontend	P1
Currency display formatting (Intl.NumberFormat)	Frontend	P1
Category override	Frontend	P2
CSV export	Backend	P3 (stretch)

11. Brand Identity & Design System

11.1 Brand tone of voice

Inspired by **Suze Orman** — one of the world's most trusted financial voices. The voice of this app is:

- **Encouraging, not condescending.** Every user is capable of financial clarity. The app assumes they're smart, just busy.
- **Direct, not preachy.** Says the honest thing once, clearly. Never repeats, never lectures.
- **Empowering, not alarming.** Frames every insight as a choice the user has power over.
- **Warm, not fluffy.** The warmth comes from respect, not from softening the truth.

Tone examples:

Avoid	Use instead
"You overspent on food this week."	"Food came in a little high this week — one fewer takeaway makes a real difference."
"Warning: high spending detected."	"Here's what I noticed — and here's one thing you can do about it."
"Great job!"	"You made a smart call cutting back on transport this week. That's how it's done."
"Error: budget exceeded."	"You've gone over on dining this month. Let's look at where to pull back."

11.2 Colour palette

Designed to convey **trust, clarity, and calm confidence** — the visual language of a premium financial tool, not a high-street bank.

Primary colours

Role	Name	Hex	Usage
Primary brand	Deep Navy	 #0F2D5C	Headings, logo, primary text on light
Action / CTA	Teal	 #0D9B8A	Buttons, active states, progress indicators
Positive / growth	Soft Emerald	 #22C997	Positive trends, "on track" states, savings milestones

Neutral palette

Role	Name	Hex	Usage
Page background	Off-white	 #F8F9FC	App background
Card surface	White	 #FFFFFF	Cards, modals, inputs
Divider	Cool Grey	 #E4E8EF	Borders, separators
Body text	Slate	 #2D3A4A	All body copy
Secondary text	Muted Grey	 #6B7A8D	Labels, captions, placeholders

Semantic colours

Role	Name	Hex	Usage
Alert / overspend	Warm Amber	 #F5A623	Spending warnings, non-essential badge
Critical	Soft Red	 #E8495A	Delete, budget exceeded (>20% over)
Essential	Steel Blue	 #4A7FC1	"Essential" badge on expenses

Colour usage rules

- Never use red for regular insights — reserve for genuinely critical states only
- Teal is the primary action colour; navy is for headings and brand only — not buttons

- Charts use a sequential palette: Teal → Emerald → Steel Blue → Amber (in that order)
 - White cards on Off-white background only — avoid grey-on-grey combinations
 - All background colours must pass WCAG AA contrast against their text colours
-

11.3 Typography

Font stack

Role	Font	Fallback	Source
Headings	DM Sans	system-ui, sans-serif	Google Fonts (free)
Body / UI	Inter	system-ui, sans-serif	Google Fonts (free)
Numeric data	Inter tabular	—	<code>font-variant-numeric: tabular-nums</code>

Why these fonts: DM Sans is modern, geometric, and confident — commanding without being aggressive, ideal for financial headings. Inter is the gold standard for UI legibility at all sizes, especially for dense financial data. Both are free, widely supported, and render crisply across all platforms.

Type scale

Element	Font	Size	Weight	Line height	Colour
Page title	DM Sans	28px	600	1.2	 #0F2D5C
Section heading	DM Sans	20px	600	1.3	 #0F2D5C
Card heading	DM Sans	16px	600	1.4	 #2D3A4A
Body text	Inter	15px	400	1.6	 #2D3A4A
Secondary / labels	Inter	13px	400	1.5	 #6B7A8D
Large numeric display	Inter	36px	600	1.1	 #0F2D5C
Small numeric / badge	Inter	12px	500	1.4	contextual
CTA button text	Inter	15px	500	1.0	 #FFFFFF
Input placeholder	Inter	15px	400	—	 #6B7A8D
AI mentor quote	DM Sans	15px	400 italic	1.7	 #2D3A4A

Typography rules

- Minimum font size: 12px (badge labels only — all reading text 13px minimum)
- Maximum two font weights per screen
- Currency amounts always use `font-variant-numeric: tabular-nums` so digits align in lists
- Category badge labels: uppercase, 12px, letter-spacing: 0.06em — the only all-caps usage in the app
- Avoid underlining text except for hyperlinks

11.4 UI component style guide

Component	Specification
Cards	White <code>(#FFFFFF)</code> , <code>(border-radius: 16px)</code> , shadow <code>(0 2px 12px rgba(15, 45, 92, 0.07))</code> , 24px internal padding
Primary button	Teal <code>(#0D9B8A)</code> fill, white text, <code>(border-radius: 10px)</code> , height 48px, Inter 15px/500

Component	Specification
Secondary button	Navy  #0F2D5C outline 1.5px, navy text, same radius and height
Ghost button	No border, Teal text, 15px/500 — for inline low-emphasis actions
Input fields	White bg,  #E4E8EF border 1px, <code>border-radius: 10px</code> , height 48px, focus ring 2px Teal
Category badges	Pill shape <code>border-radius: 999px</code> , 12px uppercase, colour bg at 12% opacity + matching full-opacity text
AI mentor card	Background  #EAF7F5, left border 3px  #0D9B8A, italic body text in DM Sans, 20px padding
Charts	Primary series in Teal; sequential palette for multiple series; labels in Slate  #2D3A4A
Spacing unit	8px base grid — use multiples: 8, 16, 24, 32, 48 only
Icon style	Outline only (e.g. Lucide or Phosphor icon sets), 20px default, Slate colour

11.5 Design principles summary

1. **Clarity first.** Every screen has one primary action. Nothing competes for attention.
2. **Numbers breathe.** Financial data needs space. Never crowd amounts or totals.
3. **Trust through restraint.** No gradients, no decorative noise, no visual complexity. Clean equals credible.
4. **Mentor, not machine.** The AI insight card should feel human — use quotation styling and conversational layout, not a system alert box.
5. **Consistent spacing.** The 8px grid is non-negotiable. Inconsistent spacing is the fastest way to make an app feel untrustworthy.

12. Localisation Specification

12.1 Language

- **v1:** English only (all UI copy, AI responses, onboarding)

- v2 roadmap: Spanish, French, Portuguese, Arabic (RTL support required)

12.2 Currency display

All amounts displayed using the user's selected currency with correct locale formatting:

Currency	Code	Example display
US Dollar	USD	\$1,240.50
Pound Sterling	GBP	£1,240.50
Euro	EUR	€1.240,50
Japanese Yen	JPY	¥124,050 (no decimals)
Indian Rupee	INR	₹1,24,050.50 (Indian grouping)
South Korean Won	KRW	₩124,050 (no decimals)

Implementation — use the native browser Intl API:

```
javascript
const formatAmount = (amount, currencyCode) =>
  new Intl.NumberFormat(undefined, {
    style: 'currency',
    currency: currencyCode,
    maximumFractionDigits: ['JPY', 'KRW'].includes(currencyCode) ? 0 : 2
  }).format(amount);
```

12.3 Date formatting

Dates display in `DD MMM YYYY` format (e.g. 07 May 2026) for cross-locale clarity. Use `Intl.DateTimeFormat` for rendering.

12.4 AI prompt localisation

All AI prompts include the user's `currency_code` and `currency_symbol` so amounts in generated text are rendered in the correct currency. No translation of prompt strings in v1 — English only.

12.5 Storing amounts

- Always store `amount` as a raw `numeric(14,2)` value — never store the formatted string
- Always store `currency_code` (ISO 4217) alongside every amount

- Format for display at render time using `Intl.NumberFormat` — never at write time

Built for Accelerator Hackathon — May 2026 · Version 1.1